

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix
from sklearn.preprocessing import LabelEncoder, StandardScaler
import warnings
warnings.filterwarnings('ignore')

plt.style.use('seaborn-v0_8')
sns.set_palette("husl")

print("✓ All libraries imported successfully!")
print("✓ Environment ready for AI Ethics project!")
```

✓ All libraries imported successfully!
✓ Environment ready for AI Ethics project!

```
In [3]: def create_synthetic_credit_dataset(n_samples=1000):
    """Create a synthetic credit dataset with potential bias patterns"""

    np.random.seed(42)
    age = np.random.normal(45, 15, n_samples)
    age = np.clip(age, 18, 80)

    age_group = pd.cut(age, bins=[0, 30, 50, 100], labels=['Young', 'Middle', 'Senior'])
    gender = np.random.choice(['Male', 'Female'], n_samples, p=[0.6, 0.4])

    base_income = np.random.lognormal(10, 0.5, n_samples)
    gender_multiplier = np.where(gender == 'Male', 1.2, 1.0)
    income = base_income * gender_multiplier

    credit_history = np.random.choice(['Good', 'Poor', 'Critical'], n_samples, p=[0.7, 0.2, 0.1])
    employment = np.random.choice(['Employed', 'Unemployed', 'Self-employed'], n_samples)
    loan_amount = np.random.lognormal(8, 0.8, n_samples)

    approval_prob = 0.3

    approval_prob += np.where(gender == 'Male', 0.15, 0.0)
    approval_prob += np.where(credit_history == 'Good', 0.3, 0.0)
    approval_prob += np.where(credit_history == 'Poor', -0.1, 0.0)
    approval_prob += np.where(credit_history == 'Critical', -0.25, 0.0)
    approval_prob += np.where(income > np.median(income), 0.2, -0.1)
    approval_prob += np.where(employment == 'Employed', 0.1, 0.0)
    approval_prob += np.where(age_group == 'Young', -0.05, 0.0)

    loan_approved = np.random.binomial(1, np.clip(approval_prob, 0, 1), n_samples)
```

```
df = pd.DataFrame({  
    'age': age,  
    'age_group': age_group,  
    'gender': gender,  
    'income': income,  
    'credit_history': credit_history,  
    'employment_status': employment,  
    'loan_amount': loan_amount,  
    'loan_approved': loan_approved  
})  
  
return df  
  
print("✓ Data generation function created")
```

✓ Data generation function created

```
In [4]: df = create_synthetic_credit_dataset(1000)  
  
print("Dataset Overview:")  
print(f"Shape: {df.shape}")  
print(f"Columns: {list(df.columns)}")  
print("\nFirst 5 rows:")  
print(df.head())  
  
print("\nDataset Info:")  
print(df.info())  
  
print("\nBasic Statistics:")  
print(df.describe())  
  
print(f"\nMissing Values:\n{df.isnull().sum()}")  
  
print(f"\nLoan Approval Distribution:")  
print(df['loan_approved'].value_counts())  
print(f"Approval Rate: {df['loan_approved'].mean():.2%}")
```

Dataset Overview:

Shape: (1000, 8)

Columns: ['age', 'age_group', 'gender', 'income', 'credit_history', 'employment_status', 'loan_amount', 'loan_approved']

First 5 rows:

	age	age_group	gender	income	credit_history	employment_status	\
0	52.450712	Senior	Male	22645.513193	Good	Employed	
1	42.926035	Middle	Male	18146.688187	Good	Employed	
2	54.715328	Senior	Female	25837.630511	Poor	Employed	
3	67.845448	Senior	Female	43054.637301	Good	Employed	
4	41.487699	Middle	Male	10349.932966	Good	Employed	

	loan_amount	loan_approved
0	447.050924	1
1	15270.157650	1
2	3775.257807	1
3	1601.667716	0
4	3380.241161	1

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1000 entries, 0 to 999

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	age	1000 non-null	float64
1	age_group	1000 non-null	category
2	gender	1000 non-null	object
3	income	1000 non-null	float64
4	credit_history	1000 non-null	object
5	employment_status	1000 non-null	object
6	loan_amount	1000 non-null	float64
7	loan_approved	1000 non-null	int32

dtypes: category(1), float64(3), int32(1), object(3)

memory usage: 52.0+ KB

None

Basic Statistics:

	age	income	loan_amount	loan_approved
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	45.376027	27825.887421	4009.046637	0.649000
std	14.193475	14711.864630	3785.066087	0.477522
min	18.000000	5840.463732	234.777193	0.000000
25%	35.286145	17412.118808	1747.633018	0.000000
50%	45.379509	24712.248052	2975.456145	1.000000
75%	54.719158	34768.977118	5050.905757	1.000000
max	80.000000	126905.224571	35965.958068	1.000000

Missing Values:

age	0
age_group	0
gender	0
income	0
credit_history	0
employment_status	0

```
loan_amount      0
loan_approved    0
dtype: int64
```

Loan Approval Distribution:

	loan_approved
1	649
0	351

Name: count, dtype: int64
 Approval Rate: 64.90%

```
In [5]: fig, axes = plt.subplots(2, 3, figsize=(18, 12))
fig.suptitle('Dataset Overview and Bias Analysis Dashboard', fontsize=16)

axes[0, 0].hist(df['age'], bins=20, alpha=0.7, color='skyblue', edgecolor='black')
axes[0, 0].set_title('Age Distribution')
axes[0, 0].set_xlabel('Age')
axes[0, 0].set_ylabel('Frequency')

gender_counts = df['gender'].value_counts()
axes[0, 1].pie(gender_counts.values, labels=gender_counts.index, autopct='%1.1f%%',
axes[0, 1].set_title('Gender Distribution')

df.boxplot(column='income', by='gender', ax=axes[0, 2])
axes[0, 2].set_title('Income Distribution by Gender')
axes[0, 2].set_xlabel('Gender')
axes[0, 2].set_ylabel('Income')

credit_counts = df['credit_history'].value_counts()
axes[1, 0].bar(credit_counts.index, credit_counts.values, color=['green', 'orange'],
axes[1, 0].set_title('Credit History Distribution')
axes[1, 0].set_xlabel('Credit History')
axes[1, 0].set_ylabel('Count')

approval_by_gender = df.groupby('gender')['loan_approved'].mean()
axes[1, 1].bar(approval_by_gender.index, approval_by_gender.values, color=['pink'],
axes[1, 1].set_title('Loan Approval Rate by Gender')
axes[1, 1].set_xlabel('Gender')
axes[1, 1].set_ylabel('Approval Rate')
axes[1, 1].set_ylim(0, 1)

for i, v in enumerate(approval_by_gender.values):
    axes[1, 1].text(i, v + 0.01, f'{v:.2%}', ha='center', va='bottom')

approval_by_age = df.groupby('age_group')['loan_approved'].mean()
axes[1, 2].bar(approval_by_age.index, approval_by_age.values, color=['lightgreen'],
axes[1, 2].set_title('Loan Approval Rate by Age Group')
axes[1, 2].set_xlabel('Age Group')
axes[1, 2].set_ylabel('Approval Rate')
axes[1, 2].set_ylim(0, 1)

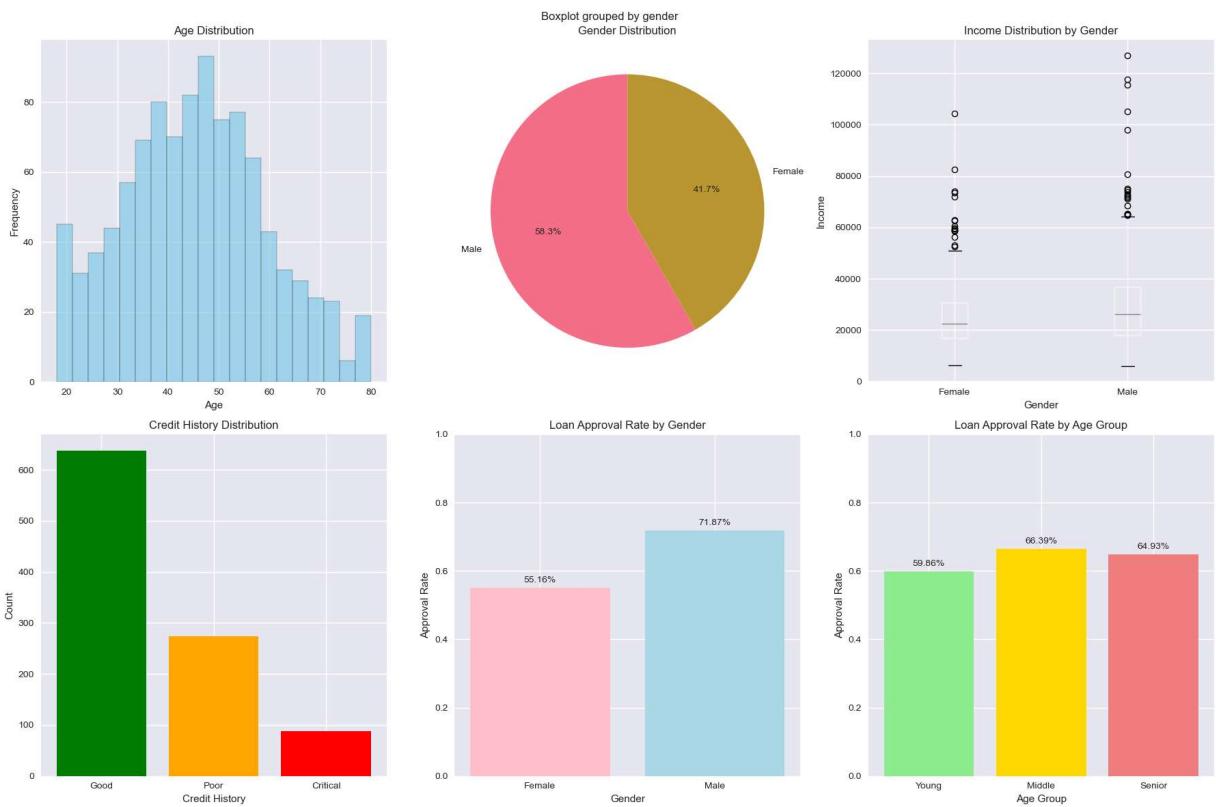
for i, v in enumerate(approval_by_age.values):
```

```

    axes[1, 2].text(i, v + 0.01, f'{v:.2%}', ha='center', va='bottom')

plt.tight_layout()
plt.show()

```



```

In [6]: class FairnessMetrics:
    """Class to calculate various fairness metrics"""

    def __init__(self, df, protected_attribute, outcome, prediction):
        self.df = df
        self.protected_attribute = protected_attribute
        self.outcome = outcome
        self.prediction = prediction

    def demographic_parity(self):
        """Calculate demographic parity difference"""
        groups = self.df[self.protected_attribute].unique()
        rates = {}

        for group in groups:
            group_data = self.df[self.df[self.protected_attribute] == group]
            rate = group_data[self.prediction].mean()
            rates[group] = rate

        max_diff = max(rates.values()) - min(rates.values())
        return rates, max_diff

    def equalized_odds_difference(self):
        """Calculate equalized odds difference"""
        groups = self.df[self.protected_attribute].unique()
        tpr_diff = {}

        for group in groups:
            group_data = self.df[self.df[self.protected_attribute] == group]
            group_outcome = group_data[self.outcome].mean()
            group_prediction = group_data[self.prediction].mean()
            tpr = group_outcome / group_prediction
            tpr_diff[group] = tpr

        max_tpr = max(tpr_diff.values())
        min_tpr = min(tpr_diff.values())
        return max_tpr - min_tpr

```

```

fpr_diff = {}

for group in groups:
    group_data = self.df[self.df[self.protected_attribute] == group]

    tp = len(group_data[(group_data[self.outcome] == 1) & (group_data[self.
p = len(group_data[group_data[self.outcome] == 1]])
tpr = tp / p if p > 0 else 0

fp = len(group_data[(group_data[self.outcome] == 0) & (group_data[self.
n = len(group_data[group_data[self.outcome] == 0]])
fpr = fp / n if n > 0 else 0

tpr_diff[group] = tpr
fpr_diff[group] = fpr

tpr_max_diff = max(tpr_diff.values()) - min(tpr_diff.values())
fpr_max_diff = max(fpr_diff.values()) - min(fpr_diff.values())

return tpr_diff, fpr_diff, tpr_max_diff, fpr_max_diff

def calibration_difference(self):
    """Calculate calibration difference"""
    groups = self.df[self.protected_attribute].unique()
    ppv_rates = {}

    for group in groups:
        group_data = self.df[self.df[self.protected_attribute] == group]

        tp = len(group_data[(group_data[self.outcome] == 1) & (group_data[self.
pred_pos = len(group_data[group_data[self.prediction] == 1]])
ppv = tp / pred_pos if pred_pos > 0 else 0

ppv_rates[group] = ppv

max_diff = max(ppv_rates.values()) - min(ppv_rates.values())
return ppv_rates, max_diff

print("✓ Fairness Metrics class created")

```

✓ Fairness Metrics class created

```

In [7]: def prepare_data(df):
    """Prepare data for machine learning"""
    df_ml = df.copy()

    le_gender = LabelEncoder()
    le_credit = LabelEncoder()
    le_employment = LabelEncoder()
    le_age_group = LabelEncoder()

    df_ml['gender_encoded'] = le_gender.fit_transform(df_ml['gender'])
    df_ml['credit_history_encoded'] = le_credit.fit_transform(df_ml['credit_history'])
    df_ml['employment_encoded'] = le_employment.fit_transform(df_ml['employment_st
    df_ml['age_group_encoded'] = le_age_group.fit_transform(df_ml['age_group'])

```

```

# Select features for model
features = ['age', 'gender_encoded', 'income', 'credit_history_encoded',
            'employment_encoded', 'loan_amount', 'age_group_encoded']

X = df_ml[features]
y = df_ml['loan_approved']

return X, y, df_ml, {
    'gender': le_gender,
    'credit': le_credit,
    'employment': le_employment,
    'age_group': le_age_group
}

X, y, df_ml, encoders = prepare_data(df)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("✓ Data prepared for machine learning")
print(f"Training set size: {X_train.shape}")
print(f"Test set size: {X_test.shape}")

```

✓ Data prepared for machine learning
 Training set size: (800, 7)
 Test set size: (200, 7)

```

In [8]: print("Training Baseline Model...")
baseline_model = RandomForestClassifier(n_estimators=100, random_state=42)
baseline_model.fit(X_train_scaled, y_train)

y_pred_baseline = baseline_model.predict(X_test_scaled)
y_prob_baseline = baseline_model.predict_proba(X_test_scaled)[:, 1]

baseline_accuracy = accuracy_score(y_test, y_pred_baseline)
baseline_precision = precision_score(y_test, y_pred_baseline)
baseline_recall = recall_score(y_test, y_pred_baseline)

print("Baseline Model Performance:")
print(f" Accuracy: {baseline_accuracy:.3f}")
print(f" Precision: {baseline_precision:.3f}")
print(f" Recall: {baseline_recall:.3f}")

df_test = df.iloc[X_test.index].copy()
df_test['prediction'] = y_pred_baseline
df_test['prediction_prob'] = y_prob_baseline

print("✓ Baseline model trained and evaluated")

```

Training Baseline Model...

Baseline Model Performance:

Accuracy: 0.710

Precision: 0.750

Recall: 0.831

✓ Baseline model trained and evaluated

```
In [9]: print("QUANTITATIVE FAIRNESS METRICS")
print("*"*50)

fairness_metrics = FairnessMetrics(df_test, 'gender', 'loan_approved', 'prediction')

dp_rates, dp_diff = fairness_metrics.demographic_parity()
print("Metric 1: Demographic Parity")
print(f" Rates: {dp_rates}")
print(f" Difference: {dp_diff:.3f}")

tpr_rates, fpr_rates, tpr_diff, fpr_diff = fairness_metrics.equalized_odds_difference()
print("\nMetric 2: Equalized Odds")
print(f" TPR Rates: {tpr_rates}")
print(f" FPR Rates: {fpr_rates}")
print(f" TPR Difference: {tpr_diff:.3f}")
print(f" FPR Difference: {fpr_diff:.3f}")

cal_rates, cal_diff = fairness_metrics.calibration_difference()
print("\nMetric 3: Calibration")
print(f" PPV Rates: {cal_rates}")
print(f" Difference: {cal_diff:.3f}")

print("\n" + "*"*50)
print("BIAS INTERPRETATION:")
print("*"*50)
print(f"• Demographic Parity Gap: {dp_diff:.1%} - {'HIGH BIAS' if dp_diff > 0.1 else 'LOW BIAS'}")
print(f"• Equalized Odds Gap: {max(tpr_diff, fpr_diff):.1%} - {'HIGH BIAS' if max(tpr_diff, fpr_diff) > 0.1 else 'LOW BIAS'}")
print(f"• Calibration Gap: {cal_diff:.1%} - {'HIGH BIAS' if cal_diff > 0.1 else 'LOW BIAS'}
```

QUANTITATIVE FAIRNESS METRICS

```
=====
Metric 1: Demographic Parity
```

```
Rates: {'Female': 0.5875, 'Male': 0.8083333333333333}
Difference: 0.221
```

```
Metric 2: Equalized Odds
```

```
TPR Rates: {'Female': 0.7291666666666666, 'Male': 0.8902439024390244}
FPR Rates: {'Female': 0.375, 'Male': 0.631578947368421}
TPR Difference: 0.161
FPR Difference: 0.257
```

```
Metric 3: Calibration
```

```
PPV Rates: {'Female': 0.7446808510638298, 'Male': 0.7525773195876289}
Difference: 0.008
```

```
=====
BIAS INTERPRETATION:
```

- Demographic Parity Gap: 22.1% - HIGH BIAS
- Equalized Odds Gap: 25.7% - HIGH BIAS
- Calibration Gap: 0.8% - LOW BIAS

```
In [17]: fig, axes = plt.subplots(2, 3, figsize=(18, 12))
fig.suptitle('Comprehensive Bias Analysis Dashboard', fontsize=16)

groups = list(dp_rates.keys())
rates = list(dp_rates.values())
colors = ['#FF6B6B', '#4CDC4']

axes[0, 0].bar(groups, rates, color=colors)
axes[0, 0].set_title('Demographic Parity by Gender')
axes[0, 0].set_ylabel('Positive Prediction Rate')
axes[0, 0].set_ylim(0, 1)

for i, v in enumerate(rates):
    axes[0, 0].text(i, v + 0.01, f'{v:.2%}', ha='center', va='bottom')

x_pos = np.arange(len(groups))
width = 0.35

tpr_values = list(tpr_rates.values())
fpr_values = list(fpr_rates.values())

axes[0, 1].bar(x_pos - width/2, tpr_values, width, label='TPR', color='#95A5A6')
axes[0, 1].bar(x_pos + width/2, fpr_values, width, label='FPR', color='#E74C3C')
axes[0, 1].set_title('Equalized Odds by Gender')
axes[0, 1].set_ylabel('Rate')
axes[0, 1].set_xticks(x_pos)
axes[0, 1].set_xticklabels(groups)
axes[0, 1].legend()
axes[0, 1].set_ylim(0, 1)

cal_values = list(cal_rates.values())
axes[0, 2].bar(groups, cal_values, color=['#9B59B6', '#F39C12'])
```

```
axes[0, 2].set_title('Calibration (PPV) by Gender')
axes[0, 2].set_ylabel('Positive Predictive Value')
axes[0, 2].set_ylim(0, 1)

for i, v in enumerate(cal_values):
    axes[0, 2].text(i, v + 0.01, f'{v:.2%}', ha='center', va='bottom')

for i, gender in enumerate(['Female', 'Male']):
    gender_data = df_test[df_test['gender'] == gender]
    cm = confusion_matrix(gender_data['loan_approved'], gender_data['prediction'])

    im = axes[1, i].imshow(cm, interpolation='nearest', cmap='Blues')
    axes[1, i].set_title(f'Confusion Matrix - {gender}')

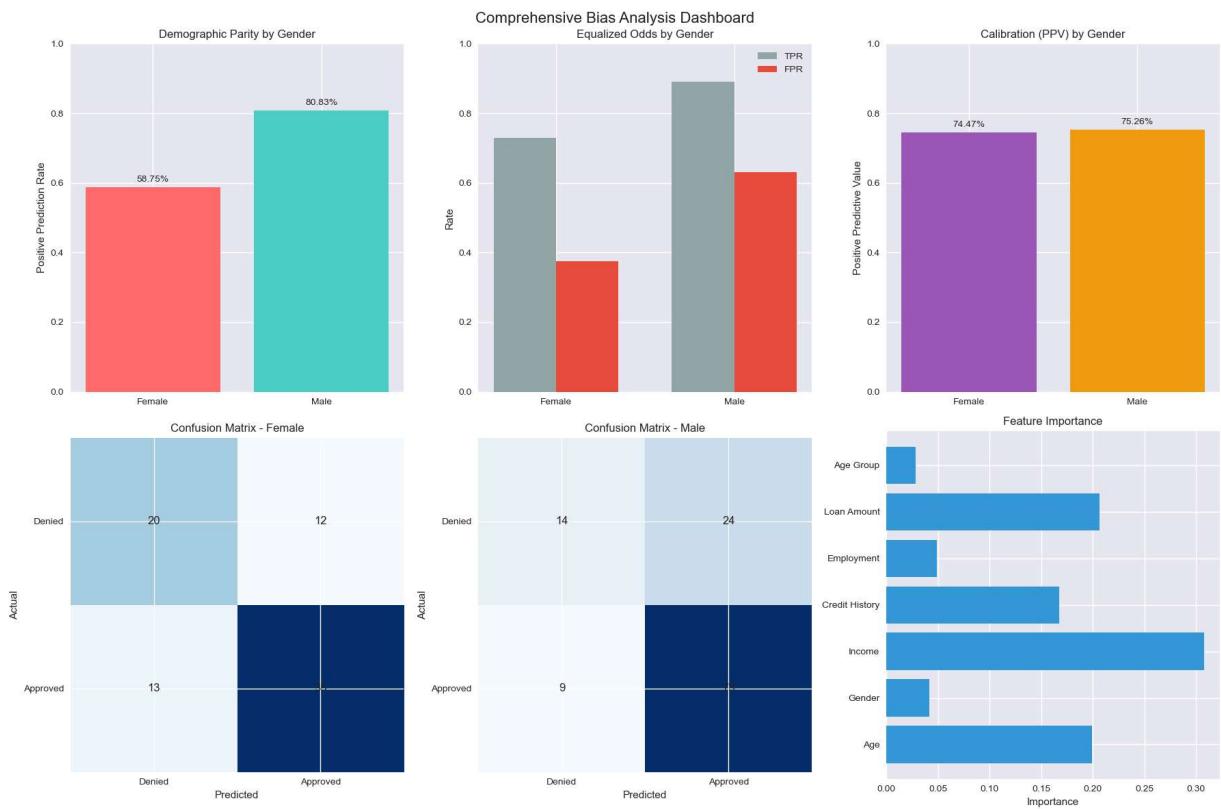
    for j in range(cm.shape[0]):
        for k in range(cm.shape[1]):
            axes[1, i].text(k, j, str(cm[j, k]), ha='center', va='center', fontsize=10)

    axes[1, i].set_xlabel('Predicted')
    axes[1, i].set_ylabel('Actual')
    axes[1, i].set_xticks([0, 1])
    axes[1, i].set_yticks([0, 1])
    axes[1, i].set_xticklabels(['Denied', 'Approved'])
    axes[1, i].set_yticklabels(['Denied', 'Approved'])

feature_importance = baseline_model.feature_importances_
feature_names = ['Age', 'Gender', 'Income', 'Credit History', 'Employment', 'Loan A

axes[1, 2].barh(feature_names, feature_importance, color="#3498DB")
axes[1, 2].set_title('Feature Importance')
axes[1, 2].set_xlabel('Importance')

plt.tight_layout()
plt.show()
```



```
In [11]: print("BIAS MITIGATION TECHNIQUE 1: REWEIGHTING")
print("*50)

def calculate_reweighting_factors(df, protected_attribute, outcome):
    """Calculate reweighting factors for bias mitigation"""

    overall_positive_rate = df[outcome].mean()

    weights = {}
    for group in df[protected_attribute].unique():
        group_data = df[df[protected_attribute] == group]
        group_positive_rate = group_data[outcome].mean()
        group_size = len(group_data)

        if group_positive_rate > 0:
            weight_positive = overall_positive_rate / group_positive_rate
        else:
            weight_positive = 1.0

        weight_negative = (1 - overall_positive_rate) / (1 - group_positive_rate)

        weights[group] = {
            'positive': weight_positive,
            'negative': weight_negative,
            'size': group_size
        }

    return weights

reweight_factors = calculate_reweighting_factors(df, 'gender', 'loan_approved')
print("Reweighting Factors:")
```

```

for group, factors in reweight_factors.items():
    print(f" {group}: Positive={factors['positive']:.3f}, Negative={factors['negative']:.3f}")

def apply_reweighting(X_train, y_train, df_train, protected_attribute, reweight_factors):
    """Apply reweighting to training data"""
    sample_weights = np.ones(len(X_train))

    for i, (idx, row) in enumerate(df_train.iterrows()):
        group = row[protected_attribute]
        outcome = row['loan_approved']

        if outcome == 1:
            sample_weights[i] = reweight_factors[group]['positive']
        else:
            sample_weights[i] = reweight_factors[group]['negative']

    return sample_weights

df_train = df.iloc[X_train.index].copy()
sample_weights = apply_reweighting(X_train, y_train, df_train, 'gender', reweight_factors)

print("Training Reweighted Model...")
reweighted_model = RandomForestClassifier(n_estimators=100, random_state=42)
reweighted_model.fit(X_train_scaled, y_train, sample_weight=sample_weights)

y_pred_reweighted = reweighted_model.predict(X_test_scaled)
df_test['prediction_reweighted'] = y_pred_reweighted

fairness_reweighted = FairnessMetrics(df_test, 'gender', 'loan_approved', 'predicted')
dp_rates_rw, dp_diff_rw = fairness_reweighted.demographic_parity()

print(f"Reweighted Model - Demographic Parity:")
print(f" Rates: {dp_rates_rw}")
print(f" Difference: {dp_diff_rw:.3f} (vs {dp_diff:.3f} baseline)")
print(f" Improvement: {((dp_diff - dp_diff_rw) / dp_diff * 100):.1f}%")

```

BIAS MITIGATION TECHNIQUE 1: REWEIGHTING

Reweighting Factors:

Male: Positive=0.903, Negative=1.248
 Female: Positive=1.177, Negative=0.783

Training Reweighted Model...

Reweighted Model - Demographic Parity:

Rates: {'Female': 0.6, 'Male': 0.7833333333333333}
 Difference: 0.183 (vs 0.221 baseline)
 Improvement: 17.0%

In [13]:

```

print("BIAS MITIGATION TECHNIQUE 2: THRESHOLD OPTIMIZATION")
print("*50")

```

```

def optimize_thresholds(df, protected_attribute, outcome, prediction_prob):
    """Optimize decision thresholds for different groups"""
    groups = df[protected_attribute].unique()
    optimal_thresholds = {}

    for group in groups:

```

```

group_data = df[df[protected_attribute] == group]

best_threshold = 0.5
best_accuracy = 0

for threshold in np.arange(0.1, 0.9, 0.05):
    pred = (group_data[prediction_prob] >= threshold).astype(int)
    accuracy = accuracy_score(group_data[outcome], pred)

    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_threshold = threshold

optimal_thresholds[group] = best_threshold

return optimal_thresholds

optimal_thresholds = optimize_thresholds(df_test, 'gender', 'loan_approved', 'predi
print("Optimal Thresholds:")
for group, threshold in optimal_thresholds.items():
    print(f" {group}: {threshold:.3f}")

def apply_optimized_thresholds(df, protected_attribute, prediction_prob, thresholds
    """Apply group-specific thresholds"""
    predictions = np.zeros(len(df))

    for group, threshold in thresholds.items():
        group_mask = df[protected_attribute] == group
        predictions[group_mask] = (df.loc[group_mask, prediction_prob] >= threshold

    return predictions

y_pred_optimized = apply_optimized_thresholds(df_test, 'gender', 'prediction_prob',
df_test['prediction_optimized'] = y_pred_optimized

fairness_optimized = FairnessMetrics(df_test, 'gender', 'loan_approved', 'predic
dp_rates_opt, dp_diff_opt = fairness_optimized.demographic_parity()

print("Optimized Thresholds - Demographic Parity:")
print(f" Rates: {dp_rates_opt}")
print(f" Difference: {dp_diff_opt:.3f} (vs {dp_diff:.3f} baseline)")
print(f" Improvement: {((dp_diff - dp_diff_opt) / dp_diff * 100):.1f}%")

```

BIAS MITIGATION TECHNIQUE 2: THRESHOLD OPTIMIZATION

=====

Optimal Thresholds:

Female: 0.350
 Male: 0.700

Optimized Thresholds - Demographic Parity:

Rates: {'Female': 0.725, 'Male': 0.675}
 Difference: 0.050 (vs 0.221 baseline)
 Improvement: 77.4%

In [14]: print("PERFORMANCE COMPARISON")
 print("=*50)

```

models = {
    'Baseline': y_pred_baseline,
    'Reweighted': y_pred_reweighted,
    'Threshold Optimized': y_pred_optimized
}

results_summary = []

for model_name, predictions in models.items():

    accuracy = accuracy_score(y_test, predictions)
    precision = precision_score(y_test, predictions)
    recall = recall_score(y_test, predictions)

    if model_name == 'Baseline':
        dp_rates_temp, dp_diff_temp = dp_rates, dp_diff
    elif model_name == 'Reweighted':
        dp_rates_temp, dp_diff_temp = dp_rates_rw, dp_diff_rw
    else:
        dp_rates_temp, dp_diff_temp = dp_rates_opt, dp_diff_opt

    results_summary.append({
        'Model': model_name,
        'Accuracy': accuracy,
        'Precision': precision,
        'Recall': recall,
        'Demographic Parity Gap': dp_diff_temp,
        'Female Approval Rate': dp_rates_temp['Female'],
        'Male Approval Rate': dp_rates_temp['Male']
    })

results_df = pd.DataFrame(results_summary)
print(results_df.round(3))

fig, axes = plt.subplots(1, 2, figsize=(15, 5))

models_list = results_df['Model'].tolist()
x_pos = np.arange(len(models_list))
width = 0.25

axes[0].bar(x_pos - width, results_df['Accuracy'], width, label='Accuracy', alpha=0.8)
axes[0].bar(x_pos, results_df['Precision'], width, label='Precision', alpha=0.8)
axes[0].bar(x_pos + width, results_df['Recall'], width, label='Recall', alpha=0.8)
axes[0].set_title('Model Performance Comparison')
axes[0].set_ylabel('Score')
axes[0].set_xticks(x_pos)
axes[0].set_xticklabels(models_list)
axes[0].legend()
axes[0].set_ylim(0, 1)

axes[1].bar(models_list, results_df['Demographic Parity Gap'], color=['red', 'orange'])
axes[1].set_title('Demographic Parity Gap Comparison')
axes[1].set_ylabel('Gap')
axes[1].set_ylim(0, max(results_df['Demographic Parity Gap']) * 1.1)

```

```

for i, v in enumerate(results_df['Demographic Parity Gap']):
    axes[1].text(i, v + 0.005, f'{v:.3f}', ha='center', va='bottom')

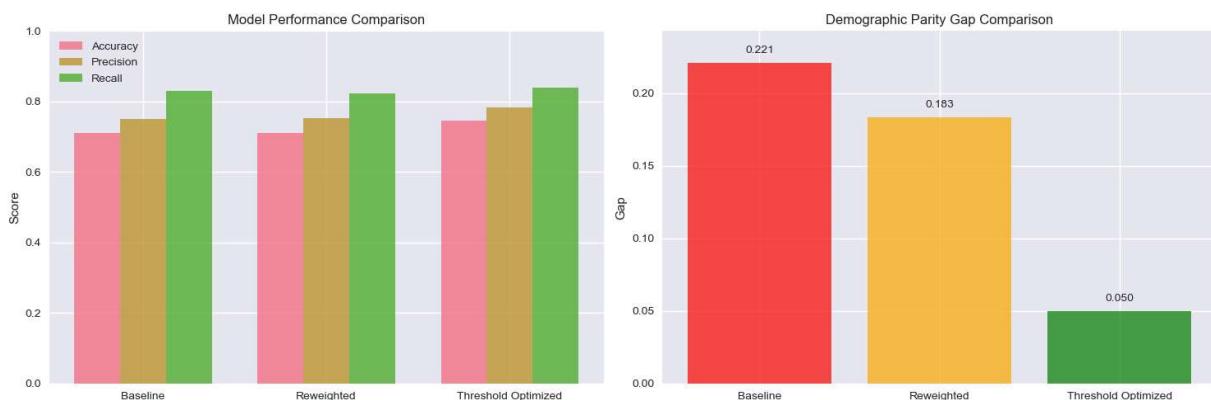
plt.tight_layout()
plt.show()

```

PERFORMANCE COMPARISON

	Model	Accuracy	Precision	Recall	Demographic Parity Gap	\
0	Baseline	0.710	0.750	0.831	0.221	
1	Reweighted	0.710	0.754	0.823	0.183	
2	Threshold Optimized	0.745	0.784	0.838	0.050	

	Female Approval Rate	Male Approval Rate
0	0.588	0.808
1	0.600	0.783
2	0.725	0.675



In [16]:

```

print("REAL-WORLD IMPLICATIONS & ETHICAL ANALYSIS")
print("*"*60)

print("1")

```

REAL-WORLD IMPLICATIONS & ETHICAL ANALYSIS

```
=====
1
```

In []: